

---

# SpikingVTG: A Spiking Detection Transformer for Video Temporal Grounding

---

Malyaban Bal<sup>1,2\*</sup>, Brian Matejek<sup>2</sup>, Susmit Jha<sup>2</sup>, Adam D. Cobb<sup>2</sup>

<sup>1</sup>The Pennsylvania State University

<sup>2</sup>Computer Science Laboratory, SRI International

## Abstract

Video Temporal Grounding (VTG) aims to retrieve precise temporal segments in a video conditioned on natural language queries. Unlike conventional neural frameworks that rely heavily on computationally expensive dense matrix multiplications, Spiking Neural Networks (SNNs)—previously underexplored in this domain—offer a unique opportunity to tackle VTG tasks through bio-plausible spike-based communication and an event-driven accumulation-based computational paradigm. We introduce SpikingVTG, a multi-modal spiking detection transformer, designed to harness the computational simplicity and sparsity of SNNs for VTG tasks. Leveraging the temporal dynamics of SNNs, our model introduces a Saliency Feedback Gating (SFG) mechanism that assigns dynamic saliency scores to video clips and applies multiplicative gating to highlight relevant clips while suppressing less informative ones. SFG enhances performance and reduces computational overhead by minimizing neural activity. We analyze the layer-wise convergence dynamics of SFG-enabled model and apply implicit differentiation at equilibrium to enable efficient, BPTT-free training. To improve generalization and maximize performance, we enable knowledge transfer by optimizing a Cos-L2 representation matching loss that aligns the layer-wise representation and attention maps of a non-spiking teacher with those of our student SpikingVTG. Additionally, we present Normalization-Free (NF)-SpikingVTG, which eliminates non-local operations like softmax and layer normalization, and an extremely quantized 1-bit (NF)-SpikingVTG variant for potential deployment on edge devices. Our models achieve competitive results on QVHighlights, Charades-STA, TACoS, and YouTube Highlights, establishing a strong baseline for multi-modal spiking VTG solutions.

## 1 Introduction

The rapid expansion of various social medias and portable smart technologies has triggered an unprecedented surge in video content. This vast influx of data has intensified the need for efficient methods to retrieve and analyze video information. Consequently, the field of Video Temporal Grounding (VTG) [1] has emerged as an important area of research. The main objective of VTG is to identify the precise segment of a video that corresponds to a given natural language query, enabling accurate and context-driven video content retrieval. In this paper, we focus on two tasks: moment retrieval [2, 3], which aims to identify video intervals relevant to a given query, and highlight detection [4], which retrieves the best candidate segment of the video in response to the query. Our work involves analyzing multimodal data—combining video content with natural language queries—to develop an effective solution to the problem. With the rise of transformer based architectures the field of VTG has seen significant advancements [5, 6]. However, these models demand substantial

---

\*Work completed while at SRI.

power and energy [7] to operate. Furthermore, VTG is inherently resource-intensive, requiring the analysis of long video sequences, leading to significant computational overhead. Additionally, applications such as **edge-based event detection** in video—e.g., identifying accidents from traffic cameras—require deploying VTG models on resource-constrained devices near the data source, in order to reduce data transfer to the cloud. These devices often operate under limited energy availability. Inspired by neural dynamics in the brain, this paper leverages bio-plausible neuronal models and learning dynamics to develop an efficient and brain-inspired solution for VTG tasks. Our model enables flexible **inference-time tradeoff** between **energy consumption and accuracy**, making it well-suited for **edge-based deployment in VTG tasks**.

**SpikingVTG Model:** We introduce SpikingVTG, a spiking detection transformer designed for efficient and accurate VTG. Built on the sparse, event-driven communication and accumulation-based computation of Spiking Neural Networks (SNNs) [8], SpikingVTG leverages the intrinsic dynamics of SNNs to offer a lightweight yet competitive alternative to conventional transformer-based approaches—eliminating the need for costly dense, real-valued matrix multiplications. The architecture comprises three key components: (i) a spiking transformer core, (ii) a Saliency Feedback Gating (SFG) mechanism, and (iii) a spiking decoder for output generation. The spiking transformer captures temporal and cross-modal dependencies, while the decoder produces task-specific outputs. The SFG mechanism addresses a critical challenge unique to VTG: given an input video composed of many clips, how can we identify those most relevant to the query? Tailored specifically for multi-modal VTG tasks, the SFG module leverages the temporal dynamics of SNNs to attend more towards the most salient segments while suppressing irrelevant clips.

**Saliency Feedback Gating Mechanism:** Operating over discrete time steps, SpikingVTG uses the intermediate spiking activity of the transformer core to dynamically estimate the relevance of each video segment. We compute a feedback-based saliency score for each segment based on the *average spiking rate* (ASR) of the output of the transformer, conditioned on the query. These scores serve as soft attention masks within a multiplicative gating mechanism, suppressing less informative segments to reduce computational overhead while enhancing focus on the most relevant candidate clips. From a neuroscience perspective, **feedback connections** are known to play a critical role in object recognition in the visual cortex [9]. Furthermore, our **feedback mechanism maintains layer-wise convergence of ASR** to equilibrium, enabling us to adopt an efficient training mechanism leveraging the equilibrium dynamics [10]. This approach circumvents the need for computationally intensive BPTT [11], and instead updates model parameters using a single backward pass, significantly improving memory efficiency.

**Cos-L2 Representation Matching (CLRM):** While our base SpikingVTG model achieves performance comparable to non-pretrained, non-spiking VTG models such as M-DETR and UniVTG [1, 6], achieving state-of-the-art results on VTG tasks typically requires transformer-based, non-spiking models like UniVTG to undergo extensive pre-training. This pre-training significantly boosts their generalization and task-specific capabilities. To enable our SpikingVTG to benefit from similar pre-training advantages—without incurring the high computational cost of training on large-scale datasets like Ego4D, Video-CC [12]—we propose a direction and scale aware CLRM loss for efficient knowledge transfer. Optimizing CLRM loss aligns the hidden states and attention score maps of a pre-trained non-spiking multi-modal transformer with the hidden state converged ASR and mean attention scores of our SpikingVTG model. By minimizing this alignment loss on the downstream task, we allow the student model to imbibe the generalization capability learnt by the pre-trained teacher without having to perform the extensive pre-training from scratch.

**Optimizations and Application to VTG Tasks:** Traditional transformer-based VTG solutions [6, 13] rely heavily on non-local normalization operations, such as softmax and layer normalization, which pose significant challenges for efficient implementation on neuromorphic hardware [14]. To address this limitation, we develop and evaluate a Normalization-Free (NF)-SpikingVTG model, which eliminates all layer normalization operations and substitutes softmax spiking attention with a ReLU and scaling-based spiking attention mechanism. Although alternative Softmax-free attention approaches have been explored in the literature [15, 16], they have primarily been applied uni-modal vision tasks. We empirically demonstrate that while these optimizations enhance computational efficiency, they result in minimal performance degradation.

To further **reduce computational complexity** and **memory footprint** [17, 18], we introduce 1-bit (NF)-SpikingVTG model, which rely primarily on **integer accumulation operations** and eliminate




Figure 1: High-level overview of the proposed SpikingVTG model. The model employs a spiking transformer core that utilizes Saliency Feedback Gating through temporal feedback connections. The model also incorporates a spiking decoder module that takes the output of the transformer core to predict parameters for the VTG task.

**all non-local operations.** This design makes it well-suited for deployment on resource-constrained edge devices. To our knowledge, this work is the first to evaluate an operational spiking detection transformer across VTG tasks, including moment retrieval and highlight detection, on datasets such as QVHighlights, Charades-STA, TACoS and Youtube Highlights. To further highlight the benefits of using SpikingVTG, we present an energy-accuracy tradeoff analysis. This demonstrates that our model is well-suited for deployment on edge devices, where its performance can be dynamically adjusted based on the available energy budget.

## 2 Related Works

**VTG Advancements:** Moment-DETR [1], a transformer encoder-decoder model introduced alongside the QVHighlights dataset, laid a strong foundation for subsequent VTG architectures. UMT [5] introduced an unified framework for solving both highlight detection and moment retrieval tasks. Due to the limited availability of trainable video data, UniVTG [6] proposed an innovative solution by unifying various VTG tasks and labels under a single formulation. This enables the development of an LLM-like pretraining framework, achieving state-of-the-art performance on VTG tasks. Although no fully spiking-based model has been explored for VTG tasks, SpikeMba [19]—primarily a non-spiking model—integrates SNN components to generate proposal sets from video data. However, since its core framework is derived from Mamba [20] and relies on floating-point matrix multiplications, SpikeMba cannot be considered a baseline for spiking models, which predominantly use accumulation-based operations. While recent VTG models have significantly improved task-specific performance, adopting a spiking framework enables us to leverage energy-accuracy tradeoff—making suitable solutions on edge devices with limited or dynamic energy supply.

**Spiking neural networks (SNNs):** SNNs have been implemented in neuromorphic systems like IBM TrueNorth [21] and Intel Loihi 2 [22], demonstrating approximately  $75\times$  greater energy efficiency compared to traditional networks running on low-power GPUs [23]. SNNs, with their energy-efficient computational framework, offer a promising solution to the resource-intensive demands of multimodal VTG tasks. While SNNs for a long time were confined in simpler vision-based tasks [24] with simple architectures, recent developments have scaled them to transformer-based models for tasks ranging from vision to language modelling [25, 26, 27], however majority of them are uni-modal and rely on non-local operations not implementable on a neuromorphic chip.

## 3 Video Temporal Grounding (VTG)

For a given video  $V$  and language query  $Q$ , we start by segmenting  $V$  into a sequence of  $L_v$  fixed-length clips, denoted as  $\{v_1, \dots, v_{L_v}\}$ . Each clip  $v_i$  has a length  $l$  and is centered at timestamp  $t_i$ . The textual query  $Q$  consists of  $L_q$  tokens, denoted as  $Q = \{q_1, \dots, q_{L_q}\}$ . Following previous studies on VTG [6], we define three parameters for each clip  $v_i = (f_i, d_i, s_i)$ , where  $f_i = 1$  if the

clip is in foreground, i.e. relevant else  $f_i = 0$ .  $d_i = [d_{s_i}, d_{e_i}] \in \mathbb{R}^2$  represent the temporal distance that converts the clip timestamp  $t_i$  to its interval boundaries. Here,  $d_i$  is valid when  $f_i = 1$ . The term  $d_{s_i}$  denotes the distance between the start of the interval and  $t_i$ , while  $d_{e_i}$  denotes the distance between the end of the interval and  $t_i$ .  $s_i \in [0, 1]$  is a continuous score that quantifies the relevance between the visual content of clip  $v_i$  and the query  $Q$ . Our proposed SpikingVTG predicts these three parameters for each video clip. In this paper, we focus on the following VTG tasks:

**Moment Retrieval:** We rank the predicted clip boundaries  $\{\tilde{b}_i\}_{i=1}^{L_v}$ , where  $b_i = [t_i - d_{s_i}, t_i + d_{e_i}]$ , based on their associated probabilities given by  $\{\tilde{f}_i\}_{i=1}^{L_v}$ . Since the predicted  $L_v$  boundaries are dense, we employ a 1-dimensional Non-Maximum Suppression (NMS) [28] with a threshold of 0.7 to eliminate highly overlapping boundary boxes, resulting in a final prediction.

**Highlight Detection** For each clip, we rank all clips based on their combined scores  $\{\tilde{f}_i + \tilde{s}_i\}_{i=1}^{L_v}$ . This combined value represents how well the clip  $i$  match with the underlying query. We then return the top clips (e.g., Top-1) as predictions.

## 4 SpikingVTG: Architecture Overview

The core computational unit of the proposed SpikingVTG model is a leaky integrate-and-fire (LIF) neuron [29]. Neurons communicate with each other using sparse, spike-based activations instead of real-valued signals, thus we can replace floating-point matrix multiplications with accumulative operations, resulting in improved computational efficiency.

### 4.1 Spiking Neural Networks

The discrete time dynamics of an LIF-based spiking neuron can be given as follows,

$$\begin{aligned} u_i[t + \delta] &= \gamma u_i[t] + W_{(i-1)}(s_{(i-1)}[t]) + b_i, \\ u_i[t + 1] &= u_i[t + \delta] - V_{th_i} s_i[t + 1], \\ a_i[t] &= \frac{\sum_{\tau=1}^t \gamma^{t-\tau} s_i[\tau]}{\sum_{\tau=1}^t \gamma^{t-\tau}}. \end{aligned} \tag{1}$$

where, at time  $t$ ,  $u_i[t]$  is the membrane potential of the  $i^{th}$  neuronal layer;  $b_i$  indicates a bias term and  $\gamma$  is the leaky term.  $W_{(i-1)}$  represents the layer-specific operation;  $t + \delta$  is an intermediate time step to determine if the neuron fired;  $V_{th_i}$  is the threshold of layer  $i$ . We use a ternary spiking model [30] in our work for spike ( $s[t + 1]$ ) generation. This improves performance while avoiding the introduction of additional floating-point multiplicative and accumulative (FP-MAC) operations. The average spiking rate (ASR)  $a_i[t]$  of neurons within each layer  $i$  at time  $t$  can be defined as a weighted-average function.

### 4.2 Spiking Transformer Core

The high-level overview of each encoder block of our spiking transformer architecture is demonstrated in Fig. 1. The model consists of  $N$  encoder layers, each consists of a spiking multi-headed attention block, followed by an intermediate layer and an output layer. Communication within and between encoder layers occurs via spikes. Furthermore, all matrix multiplications involved in linear layers and attention layer comprises of more efficient fp-accumulative (FP-ACC) operations instead of FP-MAC operations in conventional neural architectures. Detailed descriptions of each layer are provided in the Appendix A. Following architectural optimizations (Section 4.6), we replace softmax-based attention with a ReLU and scaling-based spiking attention mechanism, remove all layer normalization operations, and explore extreme quantization of linear weights.

### 4.3 Saliency Feedback Gating (SFG)

SpikingVTG operates over a specific number of convergence time steps ( $T_c$ ), with the convergence dynamics detailed in Section 4.5. This temporal processing allows us to leverage intermediate temporal outputs to dynamically update the input to the model at every time step for improved performance. This approach conforms to the feedback connections observed in the human visual

cortex [9], providing a bio-plausible explanation for its efficacy. The ASR of the final encoder layer of the Spiking Transformer core is used as a temporal feedback to compute a dynamic saliency score with the input query enabling the design of a gating mechanism (Fig. 2a), allowing selective focusing on relevant segments of the video while minimizing computation on irrelevant segments. The saliency feedback gating mechanism is shown below,

$$F_s^{v_i}[t] = \cos(\mathbf{a}_{N_v}^i[t], \mathbf{M}) := \frac{\mathbf{a}_{N_v}^i[t] \cdot \mathbf{M}}{\|\mathbf{a}_{N_v}^i[t]\|_2 \|\mathbf{M}\|_2}, \quad (2)$$

$$\bar{V}[t+1] = V[t] * \bar{F}_s^v[t],$$

$$\text{SFG}(V, \mathbf{Q}, \mathbf{a}_{N_v}) = \bar{V}[t+1] \oplus \mathbf{Q},$$

where, using attentive pooling operation, sentence representation  $\mathbf{M} = \mathbf{Q}^T \text{Softmax}(\mathbf{Q} \mathbf{W}_p)$ ,  $\mathbf{M} \in \mathbb{R}^D$ , input textual query features  $\mathbf{Q} \in \mathbb{R}^{L_q \times D}$ , input video features  $V \in \mathbb{R}^{L_v \times D}$  and  $\mathbf{W}_p \in \mathbb{R}^{D \times 1}$  is a learnable embedding and  $D$  is the hidden dimension.  $F_s^{v_i}[t]$  is the dynamic saliency score, at time  $t$ , for the  $i$ -th segment of the video. We compute  $\bar{F}_s^v$  by applying min-max normalization to  $F_s^v$ , allowing us to obtain per-clip scores within the range  $[0, 1]$ . The ASR of the output of the spiking transformer core is given as  $a_N[t] \in \mathbb{R}^{(L_v + L_q) \times D}$ .  $a_{N_v}^i[t]$  is ASR of output of the spiking transformer core, corresponding to video segment  $i$ , at time  $t$ . The output of SFG module is the concatenation of saliency feedback gated video features and query features and serves as the input to the spiking transformer core at time  $t + 1$ .

To ensure compatibility between the two modalities during cosine similarity computation, both representations are projected into a shared latent space of dimensionality  $D$ . The query representation is aggregated into a global sentence embedding via attention-based pooling with a learnable embedding vector. This adaptive mapping aligns the textual representation with the spatiotemporal semantics of the spiking video features, thereby facilitating effective cross-modal compatibility.

**SFG Computational Overhead:** The SFG layer comprises  $O(L_v \cdot D)$  floating-point multiplication operations; however, the computational overhead of this layer is significantly less than that of the transformer core which has a complexity of  $O(L^2 \cdot D + L \cdot D^2)$ , where  $L = L_v + L_q$ .

#### 4.3.1 Visualizing effect of SFG:

The dynamic saliency score ( $F_s^v$ ) achieves an equilibrium value  $F_s^*$  following the convergence of ASR of the neuronal layers of the SpikingVTG model (Fig. 3). As shown in Fig. 2b, we empirically analyze the scores per clip at equilibrium to gain insights into the functioning of the SFG based multiplicative gating mechanism. Neighboring salient clips of the target clip exhibit higher  $\bar{F}_s$  scores at equilibrium, while irrelevant clips show lower scores, highlighting the effectiveness of the SFG mechanism.

The SFG mechanism not only results in **better performance** of our SpikingVTG architecture (Table 5) but also reduces **overall neural activity** by sparsifying input spikes. Empirical results (Fig. 3b) confirm that the model with the gating mechanism exhibits a lower neural activity, particularly in the input and spiking attention layers, compared to the model without SFG.

#### 4.4 Spiking Decoder Module

The spiking decoder comprises of stacked 1-D convolutions followed by integrate-fire (IF) neuron layers ( $\gamma = 1$  in Eqn. 1), for spike generation. The spiking decoder used for predicting foreground indicator ( $f_i$ ) per clip, applies  $n_1$  1-D convolution operations with kernel size  $k_1$ , each followed by an




Figure 2: (a) Overview of the internal operations of the saliency-feedback gating mechanism. The ASR of the output of the spiking transformer core at each time step is leveraged as the feedback signal (Fig. 1). (b) Heatmap showing the scores per clip ( $\bar{F}_s$ ) at equilibrium, with the target frame for highlight detection corresponding to clip index 28.

IF layer. The final layer consists of a single output channel, and its temporal mean is passed through a sigmoid activation to produce the prediction. The spiking decoder used for  $d_i$  applies  $n_2$  1-D convolution operations with kernel size  $k_2$ , each followed by an IF layer, and the final convolution layer has two output channels to predict  $d_i = [d_{s_i}, d_{e_i}]$ , after which we compute  $b_i$ .

#### 4.5 SpikingVTG: Convergence Dynamics

The membrane potential ( $u_1$ ) of the input LIF layer to the spiking transformer core, following the SFG mechanism can be formulated as below,

$$u_1[t+1] = \gamma u_1[t] + \mathbf{SFG}(V, Q, a_{N_v}[t]) + b_1 - V_{th_1} s_1[t+1] \quad (3)$$

where,  $a^{N_v}[t]$  is the ASR of the final layer corresponding to the video features at time  $t$ ,  $V$  is the input video features,  $Q$  is the query features and  $\mathbf{SFG}$  is defined in Eqn. 2. The layer-wise convergence dynamics of the SpikingVTG with SFG is demonstrated in Fig. 3.

Following Eqn. 1, the layer-wise ASR is,  $a_i[t+1] = \frac{1}{V_{th_i}}(\hat{f}(a_{(i-1)}[t+1]) + b_i - \frac{u_i[t+1]}{\sum_{j=0}^t \gamma^j})$  where,  $\hat{f}$  is operation of layer  $i$ . Following empirical evidence (Fig. 3) and theoretical formulation [10, 26] as time  $t \rightarrow \infty$ , the layer-wise ASRs converge to equilibrium, enabling the derivation of layer-wise steady-state equations given as,

$$a_i^* = \sigma\left(\frac{1}{V_{th_i}}(\hat{f}(a_{i-1}^*) + b_i)\right) \quad (4)$$

The equation describing the steady-state ASR dynamics of the input layer is given as,

$$a_1^* = \sigma\left(\frac{1}{V_{th_1}}(\mathbf{SFG}(V, Q, a_{N_v}^*) + b_1)\right) \quad (5)$$

where, clipping function  $\sigma(x)$  clamps the values within  $[-1, 1]$ . This is because we allow ternary spikes thus ASR must be with  $[-1, 1]$ . Furthermore, the dynamic saliency score also achieves an equilibrium value  $F_s^*$  since the ASR at the final layer achieves equilibrium state  $a_{N_v}^*$ . Convergence dynamics of other layers are given in Appendix A. To analyze the overall layer-wise neural activity, which includes both positive and negative spiking event, we present the layer-wise dynamics of the absolute spiking activity events in Fig. 3b, i.e.  $act_i[t] = \frac{\sum_{i=1}^t |s_i[t]|}{t}$ .

**Training:** As described in the Section 4.4, the Spiking decoder is responsible for predicting  $\tilde{f}_i$  and  $\tilde{d}_i$  for individual video clip  $i$  and  $\tilde{s}_i$  is computed using the SFG module at equilibrium. Using these three predictions, we design a loss function that combines various components. The total loss over  $N_T$  clips in the training set is given by  $L = \frac{1}{N_T} \sum_{i=1}^{N_T} (L_{f_i} + L_{d_i} + L_{c_i})$ , where  $L_f$  is the binary cross-entropy loss for the indicator variable  $f_i$ ,  $L_d$  combines smooth L1 loss with IoU loss [31] for the predicted boundaries, and  $L_c$  is an optional loss incorporating intra- and inter-video contrastive learning [32]. Detailed formulation of the loss functions is in Appendix B.

During training, leveraging implicit differentiation [33] at equilibrium, only ASR values at equilibrium are used,  $\frac{\partial L(a^*)}{\partial \theta} = -\frac{\partial L(a^*)}{\partial a^*} (J_{g_\theta}^{-1}|_{a^*}) \frac{\partial f_\theta(a^*)}{\partial \theta}$  where,  $\theta$  is the model parameters,  $g_\theta(a) = f_\theta(a) - a$ ,



Figure 3: Results from a random QVHighlights input passed through SpikingVTG models. (a) Layer-wise mean ASR convergence over operating time steps for a sample spiking transformer encoder layer (Fig. 1); note ASR can be negative due to ternary spikes. (b) Layer-wise mean spiking activity ( $act_i[t]$ , averaged over neurons) versus time steps. Model with SFG exhibit reduced activity in both input and spiking attention layers, highlighting SFG’s role in minimizing neural activity.

$f$  is the steady-state equation of ASR,  $J^{-1}$  is the inverse Jacobian of  $g_\theta$  when  $a = a^*$ , i.e., at equilibrium. Thus, unlike BPTT, we do not need to store the intermediate computational graph and the model parameters can be updated using a single backpropagation step.

#### 4.5.1 Cos-L2 Representation Matching (CLRM)

We train the base SpikingVTG model on VTG tasks such as QVHighlights [1], achieving comparable performance to a similarly scaled non-spiking transformer-based model, UniVTG [6], as demonstrated in Table 1. The later to improve its performance further uses pre-training on large datasets such as Video-CC and Ego4D. Since, extensive pretraining of the SpikingVTG is considerably resource intensive we propose an effective knowledge transfer strategy. This mechanism enables SpikingVTG to inherit the generalization capabilities learned by the pre-trained detection transformer models like UniVTG model. Importantly, since the primary benefit of spiking architectures manifests during inference on resource-constrained edge devices, this knowledge transfer is a one-time process, which significantly improves performance.

**Hidden State Matching Loss:** To align the hidden state (output of individual encoder layer) of the non-spiking multi-modal transformer model with the converged ASR (4.5) of the corresponding layer of the SpikingVTG model, we propose a hybrid loss function combining a squared cosine similarity based directional loss and a L2-norm based loss for minimizing the scale difference. For each layer, the student representation is projected into the feature space of the teacher via a learnable linear transformation  $W_d \in \mathbb{R}^{d_s \times d_t}$ . The total loss is:

$$\mathcal{L}_{\text{rep}} = \frac{1}{B \times L} \sum_{i=1}^N \sum_{j=1}^B \sum_{k=1}^L \left[ \lambda_{\cos} \cdot \left( 1 - \cos \left( \theta_{i,j,k}^{\text{rep}} \right) \right)^2 + \lambda_{l_2} \cdot \left\| \mathbf{s}_i^{(j,k)} - \mathbf{t}_i^{(j,k)} \right\|_2^2 \right] \quad (6)$$

where  $N$  is total number of encoder layers,  $B$  is batch size,  $L$  is length of sequence,  $\lambda_{\cos}$ ,  $\lambda_{l_2}$  are hyperparameters and the cosine similarity is computed as:  $\cos(\theta_{i,j,k}^{\text{rep}}) = \frac{\mathbf{s}_i^{(j,k)} \cdot \mathbf{t}_i^{(j,k)}}{\|\mathbf{s}_i^{(j,k)}\|_2 \cdot \|\mathbf{t}_i^{(j,k)}\|_2}$ , where,  $\mathbf{s}_i^{(j,k)} = a_{r_i}^{*(j,k)} W_d \in \mathbb{R}^{d_t}$  is the projected student representation and  $\mathbf{t}_i^{(j,k)} = T_{r_i}^{(j,k)} \in \mathbb{R}^{d_t}$  is the teacher representation at layer  $i$ , sequence position  $k$ , and batch index  $j$ .  $a_{r_i}^* \in \mathbb{R}^{B \times L \times d_s}$  and  $T_{r_i} \in \mathbb{R}^{B \times L \times d_t}$  are the pre-activation outputs from student and teacher, respectively.

**Attention Score Matching Loss:** We align the attention score map ( $\mathbf{A}_i^T \in \mathbb{R}^{B \times L \times L}$ ) of encoder layer  $i$  of the non-spiking model with the mean attention score ( $\mathbf{A}_i^S = \frac{1}{T_c} \sum_{t=1}^{T_c} A_i^S[t] \in \mathbb{R}^{B \times L \times L}$ ) of our corresponding converged SpikingVTG model. The total attention map alignment loss ( $\mathcal{L}_{\text{att}}$ ) is then computed following Eqn. 6. This loss term encourages the student to learn cross-modal attention behavior consistent with the pre-trained teacher.

The cumulative representation matching loss is given as  $\mathcal{L}_{\text{CLRM}} = \mathcal{L}_{\text{rep}} + \mathcal{L}_{\text{att}}$ . Optimizing this loss enables our SpikingVTG model to learn the generalability enabling better performance as demonstrated in Table 1.

#### 4.6 Normalization-Free and Quantized SpikingVTG

We perform two key optimizations: removal of non-local operations and 1-bit weight quantization. Although SpikingVTG replaces floating-point MACs with ternary spikes, it retains softmax and layer normalization, which are inefficient for resource-constrained hardware. We eliminate these by introducing a Normalization-Free (NF) SpikingVTG.

In NF-SpikingVTG, softmax in the spiking attention is replaced with a ReLU followed by scaling with  $1/L$ , reducing compute overhead. Given  $d$ -dimensional queries, keys, and values  $\{q_i[t], s_{k_i}[t], s_{v_i}[t]\}_{i=1}^L$ , at time  $t$ , the attention weights  $\alpha_{ij}$  are computed as follows:

$$\alpha_{ij}[t] = \phi \left( [q_i[t]^\top s_{k_1}[t], \dots, q_i[t]^\top s_{k_L}[t]] \right)_j \quad (7)$$

where,  $\phi$  is a custom kernel consisting of ReLU operation followed by scaling with  $L^{-1}$ . We further remove all layer normalization layers to eliminate additional non-local operations.

To reduce memory and computational cost, we introduce 1-bit NF-SpikingVTG by binarizing weights. Each weight matrix  $W \in \mathbb{R}^{n \times m}$  is zero-centered and quantized as  $W_q = \text{sgn}\left(W - \frac{1}{nm} \sum_{i,j} W_{ij}\right)$ ,  $\beta = \frac{1}{nm} \sum_{i,j} |W_{ij}|$ , where  $W_q \in \{-1, +1\}$ . The output of each quantized linear layer is scaled by  $\beta$ , yielding a binary-weight, ternary-activation model that supports efficient integer-ACC operations. 1-bit NF-SpikingVTG achieves competitive performance while drastically reducing memory and compute, making it well-suited for edge deployment.

Table 2: Performance comparison of our SpikingVTG model against non-spiking VTG solutions on **test sets** of QVHighlights-MR and Charades-STA datasets for moment retrieval task.

Method	SNN	QVHighlights-MR				Charades-STA			
		@0.5	@0.7	mAP@0.5	mAP@0.75	@0.3	@0.5	@0.7	mIoU
M-DETR [1]	No	52.89	33.02	54.82	29.40	65.83	52.07	30.59	45.54
UMT [5]	No	56.23	41.18	53.83	37.01	-	49.35	26.16	-
UniVTG [6]	No	58.86	40.86	57.60	35.59	70.81	58.01	35.65	50.10
UniVTG+PT [6]	No	65.43	50.06	64.06	45.02	72.63	60.19	38.55	52.17
UVCOM [34]	No	63.55	47.47	63.37	42.67	-	56.69	34.76	-
SpikeMba [19]	No	64.13	49.42	-	43.67	71.24	59.65	36.12	51.74
BAM-DETR [35]	No	62.71	48.64	64.57	46.33	72.93	59.95	39.38	52.33
TR-DETR [36]	No	64.66	48.96	63.98	43.73	-	57.61	33.52	-
CG-DETR [37]	No	65.43	48.38	64.51	42.77	70.40	58.40	36.30	50.10
LLMEPET [38]	No	66.73	49.94	65.76	43.91	70.91	-	36.49	50.25
<b>SpikingVTG</b>	Yes	<b>65.29</b>	<b>48.18</b>	<b>64.31</b>	<b>42.25</b>	<b>71.20</b>	<b>58.73</b>	<b>37.16</b>	<b>50.62</b>

## 5 Experimentation

We evaluate SpikingVTG variants on moment retrieval and highlight detection tasks using the QVHighlights, Charades-STA, TACoS and Youtube Highlight datasets. Since, to the best of our knowledge, our proposed model is the first spiking detection transformer evaluated on VTG tasks, we compare its performance against sota non-spiking detection transformers. Additional training, dataset, evaluation metric, hyper-parameter and experimental details are provided in Appendix C & E. The experiments were run on a NVIDIA RTX A6000 GPU with 48GB memory.

### 5.1 Results

SpikingVTG establishes a baseline for spiking models on VTG tasks. The results are shown in Table 2, 3 & 4. Our model achieves competitive results compared to the current SOTA non-spiking models.

**Training & Inference Metrics:** In the CLRM-based knowledge transfer stage, the memory requirement is 20GB when using a batch size of 32 on the QVHighlights dataset. In contrast, replacing our training method with BPTT would require over 100GB of memory for  $T_c > 10$ , making BPTT computationally infeasible. Training on true labels with similar batch size requires 8GB of

Table 3: Performance comparison on the **test set** of TACoS for **moment retrieval task**.

Method	SNN	TACoS			
		@0.3	@0.5	@0.7	mIoU
M-DETR [1]	No	37.97	24.67	11.97	25.49
2D-TAN [6]	No	40.01	27.99	12.92	27.22
UniVTG [6]	No	51.44	34.97	17.35	33.60
QD-DETR [39]	No	52.39	36.77	21.07	35.76
CG-DETR [37]	No	52.23	-	22.23	36.48
UniVTG+PT [6]	No	56.11	43.44	24.27	38.63
SpikeMba [19]	No	51.98	39.34	22.83	35.81
<b>SpikingVTG</b>	Yes	<b>54.71</b>	<b>39.27</b>	<b>21.84</b>	<b>36.02</b>

Table 4: Performance comparison of our SpikingVTG model against other non-spiking VTG solutions on the **test set** of the QVHighlights and Youtube Highlights for **highlight detection task**.

Method	SNN	QVHighlights-HD		Youtube Highlights					
		mAP	HIT@1	Dog	Gym.	Skating	Skiing	Parking	Surfing
UMT [5]	No	38.18	59.99	65.9	75.2	71.8	72.3	81.6	82.7
UniVTG [6]	No	38.20	60.96	71.8	76.5	73.3	73.2	73.9	82.2
UniVTG+PT [6]	No	40.54	66.28	74.3	79.0	84.9	75.1	74.4	83.9
QD-DETR [39]	No	38.90	62.40	72.2	77.4	72.7	72.8	71.0	80.6
SpikeMba [5]	No	-	-	74.4	75.4	74.3	75.5	-	75.5
CG-DETR [37]	No	40.30	66.20	76.3	76.1	76.0	75.1	70.0	81.9
UVCOM [34]	No	39.98	65.58	73.8	77.1	76.0	75.1	75.7	82.7
LLMEPET [38]	No	38.18	59.99	73.6	73.2	75.3	74.0	72.5	82.5
<b>SpikingVTG</b>	Yes	<b>40.46</b>	<b>65.82</b>	<b>73.9</b>	<b>78.1</b>	<b>80.1</b>	<b>74.2</b>	<b>72.2</b>	<b>81.7</b>

memory. The clock time for 50 epochs of training on QVHighlights is  $\approx 5$  HOURS. Inference latency on entire test set of QVHighlights ranges from 25 sec ( $T_c = 2$ ) to  $\approx 4$  mins ( $T_c = 20$ ).

Table 5: **Ablation Study** of SpikingVTG variants on the evaluation set of QVHighlights.

Method	QVHighlights-MR		QVHighlights-HD		Operations	Local	Activity	Energy
	@0.5	@0.7	mAP	HIT@1				
Pre-trained UniVTG (sota)	67.35	52.65	41.34	68.77	FP-MAC	$\times$	1.0	23.92mJ
SpikingVTG <b>without SFG</b>	64.94	47.21	40.49	67.37	FP-ACC	$\times$	0.41	15.2mJ
SpikingVTG <b>with SFG</b>	<b>67.58</b>	<b>50.82</b>	<b>40.81</b>	<b>68.64</b>	FP-ACC	$\times$	0.34	13.8mJ
(NF)-SpikingVTG w/ SFG	66.59	48.31	40.61	67.73	FP-ACC	$\checkmark$	0.25	10.1mJ
1-bit (NF)-SpikingVTG w/ SFG	65.31	47.48	40.35	67.30	INT-ACC	$\checkmark$	<b>0.19</b>	<b>1.3mJ</b>
1-bit (NF)-SpikingVTG w/ ReLU	65.91	47.04	40.16	67.07	INT-ACC	$\checkmark$	<b>0.19</b>	<b>1.3mJ</b>

**Ablation Study:** As demonstrated in Table 5, the inclusion of the SFG mechanism enhances performance compared to the model without SFG. It results in reduced neuronal activity and overall less energy consumption. Thus, we enable SFG for the SpikingVTG variants we discuss next. Without non-local operations, the (NF)-SpikingVTG model achieves competitive performance even compared to other SOTA non-spiking VTG models. Although the 1-bit (NF)-SpikingVTG variant shows a slight reduction in performance, it is highly memory efficient and involves simpler INT-ACC computations, resulting in order of magnitude less energy consumption. Along with energy consumption ( $T_c = 10$ ) Table 5, also highlights the **sparsity** in the SpikingVTG variants particularly in the 1-bit version underscoring considerably **reduced mean neural activity**. For a more hardware friendly model, we train a variant of 1-bit (NF)-SpikingVTG by replacing all GELU layers with ReLU layer [40] and observe minimal degradation in performance. Extensive hyper-parameter study is done in Appendix E.

**Analysis of Energy Efficiency:** We analyze the test-time energy efficiency of SpikingVTG variants compared to a non-spiking transformer-based model with same depth and hidden dimensions. The analysis considers arithmetic operation costs using a 45nm CMOS technology with 32-bit precision, where FP-MAC, FP-ACC, and INT-ACC ops. consume 4.6 pJ, 0.9 pJ, and 0.1 pJ respectively [41]. The energy cost for a non-spiking transformer encoder layer based on total FLOPs used in attention and linear layers, is  $E_A = [(3LD^2) + (LD^2 + L^2D) + (2LD^3)] \times 4.6$  pJ. Considering  $L = 200$ ,  $D = 1024$ , we compute  $E_A = 5.98$  mJ. Since, UniVTG has 4 encoder layers, so total energy is 23.92mJ.

In contrast, the 1-bit (NF)-SpikingVTG operates over  $T_c$  time steps. Its per time-step cost is:  $E_{S_t} = [(3 \cdot IFR_{in} \cdot LD^2) + (IFR_k \cdot LD^2 + IFR_v \cdot L^2D) + (IFR_{attn} \cdot LD^2) + (IFR_{int} \cdot LD^2)] \times 0.1$  pJ, where  $IFR_l$  denotes the mean activity of component  $l$ . For our 1-bit (NF)-SpikingVTG model, which uses INT-ACC operations, we empirically measure:  $IFR_{in} = 0.40$ ,  $IFR_k = 0.18$ ,  $IFR_v = 0.19$ ,  $IFR_{attn} = 0.03$ , and  $IFR_{int} = 0.09$ , resulting in  $E_{S_t} = 0.03$  mJ. With  $T_c = 10$ , the total spiking energy is  $E_S = T_c \cdot E_{S_t} = 0.3$  mJ, yielding an energy efficiency of  $E_f = E_A / E_S = 19.93$ . Model-wise energy comparison is shown in Table 5.

**Energy-Accuracy Tradeoff:** Unlike conventional non-spiking VTG solutions, SpikingVTG enables a controllable trade-off between performance and energy consumption, as total energy usage scales with the number of operating time steps. In Fig. 4, we demonstrate this performance-energy tradeoff achieved using 1-bit (NF)-SpikingVTG model by running it for different number of time steps. Notably, when running the model for just 2 time steps, SpikingVTG attains a HIT@1 score of 66.4 (compared to the SOTA score of 68.77), while achieving an energy efficiency factor  $E_f \approx 100$ .




Figure 4: Graph showing tradeoff of energy efficiency ( $E_f$ ) and HIT@1 on QVHighlights-HD for varying time steps.

## 6 Conclusions

In this paper we propose SpikingVTG, a bio-inspired computationally efficient solution for VTG tasks in resource constraint environment. We propose a saliency feedback gating mechanism, which leverages the temporal dynamics of the model to improve performance while lowering computational costs by reducing overall neuronal activity across the model. We empirically demonstrate the convergence dynamics of the SpikingVTG model and leverage the formulated steady-state equations to efficiently train our model using implicit differentiation at equilibrium. To improve generalization, we propose a Cos-L2 Representation Matching loss, enabling knowledge transfer from non-spiking VTG models and yielding substantial performance gains. We further optimize the model by removing non-local operations and applying extreme quantization, leading to the 1-bit NF-SpikingVTG. Our framework significantly improves computational efficiency and model compactness, enabling tunable energy-accuracy tradeoffs on low-resource devices. While this work presents the first spiking solution for VTG with competitive performance, there remains room for improvement in fully closing the performance gap. Future research can build upon our framework to further advance spiking models in this domain.

## 7 Acknowledgments

This work was supported in part by the United States Air Force and Defense Advanced Research Projects Agency (DARPA) under Contract No. FA8750-23-C-0519 and the U.S. Army Research Laboratory Cooperative Research Agreement W911NF-17-2-0196. Any opinions, findings and conclusions or recommendations expressed in this material are those of the authors and do not necessarily reflect the Department of Defense or the United States Government.

## References

- [1] Jie Lei, Tamara L Berg, and Mohit Bansal. Detecting moments and highlights in videos via natural language queries. *Advances in Neural Information Processing Systems*, 34:11846–11858, 2021.
- [2] Songyang Zhang, Houwen Peng, Jianlong Fu, and Jiebo Luo. Learning 2d temporal adjacent networks for moment localization with natural language. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 34, pages 12870–12877, 2020.
- [3] Jonghwan Mun, Minsu Cho, and Bohyung Han. Local-global video-text interactions for temporal grounding. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 10810–10819, 2020.
- [4] Fa-Ting Hong, Xuanteng Huang, Wei-Hong Li, and Wei-Shi Zheng. Mini-net: Multiple instance ranking network for video highlight detection. In *Computer Vision–ECCV 2020: 16th European Conference, Glasgow, UK, August 23–28, 2020, Proceedings, Part XIII 16*, pages 345–360. Springer, 2020.
- [5] Ye Liu, Siyuan Li, Yang Wu, Chang-Wen Chen, Ying Shan, and Xiaohu Qie. Umt: Unified multi-modal transformers for joint video moment retrieval and highlight detection. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 3042–3051, 2022.
- [6] Kevin Qinghong Lin, Pengchuan Zhang, Joya Chen, Shraman Pramanick, Difei Gao, Alex Jin-peng Wang, Rui Yan, and Mike Zheng Shou. Univtg: Towards unified video-language temporal grounding. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 2794–2804, 2023.
- [7] Siddharth Samsi, Dan Zhao, Joseph McDonald, Baolin Li, Adam Michaleas, Michael Jones, William Bergeron, Jeremy Kepner, Devesh Tiwari, and Vijay Gadepally. From words to watts: Benchmarking the energy costs of large language model inference. In *2023 IEEE High Performance Extreme Computing Conference (HPEC)*, pages 1–9. IEEE, 2023.
- [8] Samanwoy Ghosh-Dastidar and Hojjat Adeli. Spiking neural networks. *International journal of neural systems*, 19(04):295–308, 2009.

- [9] Kohitij Kar, Jonas Kubilius, Kailyn Schmidt, Elias B Issa, and James J DiCarlo. Evidence that recurrent circuits are critical to the ventral stream’s execution of core object recognition behavior. *Nature neuroscience*, 22(6):974–983, 2019.
- [10] Mingqing Xiao, Qingyan Meng, Zongpeng Zhang, Yisen Wang, and Zhouchen Lin. Training feedback spiking neural networks by implicit differentiation on the equilibrium state. *Advances in Neural Information Processing Systems*, 34:14516–14528, 2021.
- [11] Emre O Neftci, Hesham Mostafa, and Friedemann Zenke. Surrogate gradient learning in spiking neural networks: Bringing the power of gradient-based optimization to spiking neural networks. *IEEE Signal Processing Magazine*, 36(6):51–63, 2019.
- [12] Arsha Nagrani, Paul Hongseok Seo, Bryan Seybold, Anja Hauth, Santiago Manen, Chen Sun, and Cordelia Schmid. Learning audio-video modalities from image captions. In *European Conference on Computer Vision*, pages 407–426. Springer, 2022.
- [13] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need, 2017.
- [14] Amar Shrestha, Haowen Fang, Zaidao Mei, Daniel Patrick Rider, Qing Wu, and Qinru Qiu. A survey on neuromorphic computing: Models and hardware. *IEEE Circuits and Systems Magazine*, 22(2):6–35, 2022.
- [15] Soroush Abbasi Koohpayegani and Hamed Pirsiavash. Sima: Simple softmax-free attention for vision transformers. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, pages 2607–2617, 2024.
- [16] Boxun Xu, Hejia Geng, Yuxuan Yin, and Peng Li. Ds2ta: Denoising spiking transformer with attenuated spatiotemporal attention. *arXiv preprint arXiv:2409.15375*, 2024.
- [17] Hongyu Wang, Shuming Ma, Li Dong, Shaohan Huang, Huaijie Wang, Lingxiao Ma, Fan Yang, Ruijing Wang, Yi Wu, and Furu Wei. Bitnet: Scaling 1-bit transformers for large language models. *arXiv preprint arXiv:2310.11453*, 2023.
- [18] Malyaban Bal, Yi Jiang, and Abhroneil Sengupta. Exploring extreme quantization in spiking language models. In *2024 International Conference on Neuromorphic Systems (ICONS)*, pages 272–276. IEEE, 2024.
- [19] Wenrui Li, Xiaopeng Hong, and Xiaopeng Fan. Spikemba: Multi-modal spiking saliency mamba for temporal video grounding. *arXiv preprint arXiv:2404.01174*, 2024.
- [20] Albert Gu and Tri Dao. Mamba: Linear-time sequence modeling with selective state spaces. *arXiv preprint arXiv:2312.00752*, 2023.
- [21] Michael V DeBole, Brian Taba, Arnon Amir, Philipp Akopyan, Alexander Andreopoulos, William P Risk, Jeff Kusnitz, Carlos Ortega Otero, Tapan K Nayak, Rathinakumar Appuswamy, et al. Truenorth: Accelerating from zero to 64 million neurons in 10 years. *Computer*, 52(5):20–29, 2019.
- [22] Mike Davies, Andreas Wild, Garrick Orchard, Yulia Sandamirskaya, Gabriel A Fonseca Guerra, Prasad Joshi, Philipp Plank, and Sumedh R Risbud. Advancing neuromorphic computing with loih: A survey of results and outlook. *Proceedings of the IEEE*, 109(5):911–934, 2021.
- [23] Guangzhi Tang, Neelesh Kumar, and Konstantinos P. Michmizos. Reinforcement co-learning of deep and spiking neural networks for energy-efficient mapless navigation with neuromorphic hardware. In *2020 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 6090–6097, 2020.
- [24] Kashu Yamazaki, Viet-Khoa Vo-Ho, Darshan Bulsara, and Ngan Le. Spiking neural networks and their applications: A review. *Brain Sciences*, 12(7):863, 2022.
- [25] Zhaokun Zhou, Yuesheng Zhu, Chao He, Yaowei Wang, Shuicheng Yan, Yonghong Tian, and Li Yuan. Spikformer: When spiking neural network meets transformer. *arXiv preprint arXiv:2209.15425*, 2022.

- [26] Malyaban Bal and Abhroni Sengupta. Spikingbert: Distilling bert to train spiking language models using implicit differentiation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, pages 10998–11006, 2024.
- [27] Rui-Jie Zhu, Qihang Zhao, and Jason K Eshraghian. Spikegpt: Generative pre-trained language model with spiking neural networks. *arXiv preprint arXiv:2302.13939*, 2023.
- [28] Jan Hosang, Rodrigo Benenson, and Bernt Schiele. Learning non-maximum suppression. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 4507–4515, 2017.
- [29] Sangya Dutta, Vinay Kumar, Aditya Shukla, Nihar R Mohapatra, and Udayan Ganguly. Leaky integrate and fire neuron by charge-discharge dynamics in floating-body mosfet. *Scientific reports*, 7(1):8257, 2017.
- [30] Yufei Guo, Yuanpei Chen, Xiaode Liu, Weihang Peng, Yuhan Zhang, Xuhui Huang, and Zhe Ma. Ternary spike: Learning ternary spikes for spiking neural networks. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, pages 12244–12252, 2024.
- [31] Hamid Rezatofighi, Nathan Tsoi, JunYoung Gwak, Amir Sadeghian, Ian Reid, and Silvio Savarese. Generalized intersection over union: A metric and a loss for bounding box regression. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 658–666, 2019.
- [32] Ting Chen, Simon Kornblith, Mohammad Norouzi, and Geoffrey Hinton. A simple framework for contrastive learning of visual representations. In *International conference on machine learning*, pages 1597–1607. PMLR, 2020.
- [33] Shaojie Bai, J Zico Kolter, and Vladlen Koltun. Deep equilibrium models. *Advances in Neural Information Processing Systems*, 32, 2019.
- [34] Yicheng Xiao, Zhuoyan Luo, Yong Liu, Yue Ma, Hengwei Bian, Yatai Ji, Yujiu Yang, and Xiu Li. Bridging the gap: A unified video comprehension framework for moment retrieval and highlight detection, 2023.
- [35] Pilhyeon Lee and Hyeran Byun. Bam-detr: Boundary-aligned moment detection transformer for temporal sentence grounding in videos, 2024.
- [36] Hao Sun, Mingyao Zhou, Wenjing Chen, and Wei Xie. Tr-detr: Task-reciprocal transformer for joint moment retrieval and highlight detection, 2024.
- [37] WonJun Moon, Sangeek Hyun, SuBeen Lee, and Jae-Pil Heo. Correlation-guided query-dependency calibration for video temporal grounding, 2024.
- [38] Yiyang Jiang, Wengyu Zhang, Xulu Zhang, Xiao-Yong Wei, Chang Wen Chen, and Qing Li. Prior knowledge integration via llm encoding and pseudo event regulation for video moment retrieval. In *Proceedings of the 32nd ACM International Conference on Multimedia*, MM ’24, page 7249–7258. ACM, October 2024.
- [39] WonJun Moon, Sangeek Hyun, SangUk Park, Dongchan Park, and Jae-Pil Heo. Query-dependent video representation for moment retrieval and highlight detection. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 23023–23033, 2023.
- [40] Jonathan Timcheck, Sumit Bam Shrestha, Daniel Ben Dayan Rubin, Adam Kupryjanow, Garrick Orchard, Lukasz Pindor, Timothy Shea, and Mike Davies. The intel neuromorphic dns challenge. *Neuromorphic Computing and Engineering*, 3(3):034005, 2023.
- [41] Song Han, Jeff Pool, John Tran, and William J. Dally. Learning both weights and connections for efficient neural networks, 2015.
- [42] Jiyang Gao, Chen Sun, Zhenheng Yang, and Ram Nevatia. Tall: Temporal activity localization via language query. In *Proceedings of the IEEE international conference on computer vision*, pages 5267–5275, 2017.

- [43] Michaela Regneri, Marcus Rohrbach, Dominikus Wetzel, Stefan Thater, Bernt Schiele, and Manfred Pinkal. Grounding action descriptions in videos. *Transactions of the Association for Computational Linguistics*, 1:25–36, 2013.
- [44] Min Sun, Ali Farhadi, and Steve Seitz. Ranking domain-specific highlights by analyzing edited videos. In *Computer Vision–ECCV 2014: 13th European Conference, Zurich, Switzerland, September 6–12, 2014, Proceedings, Part I 13*, pages 787–802. Springer, 2014.